

Double Distributionally Robust Bid Shading for First Price Auctions

Yanlin Qu

Columbia Business School

qu.yanlin@columbia.edu

Joint work with Ravi Kant, Yan Chen, Brendan
Kitts, San Gultekin, Aaron Flores, Jose Blanchet

yahoo! dsp

Working of a Demand Side Platform



VERTOZ

First-Price Auction

Profit from shading



Second-Price Auction

Always bid truthfully



Overview

1. Motivation

- Noisy real-time bidding system
- Distributionally robust optimization

2. Distributionally Robust Bid Shading

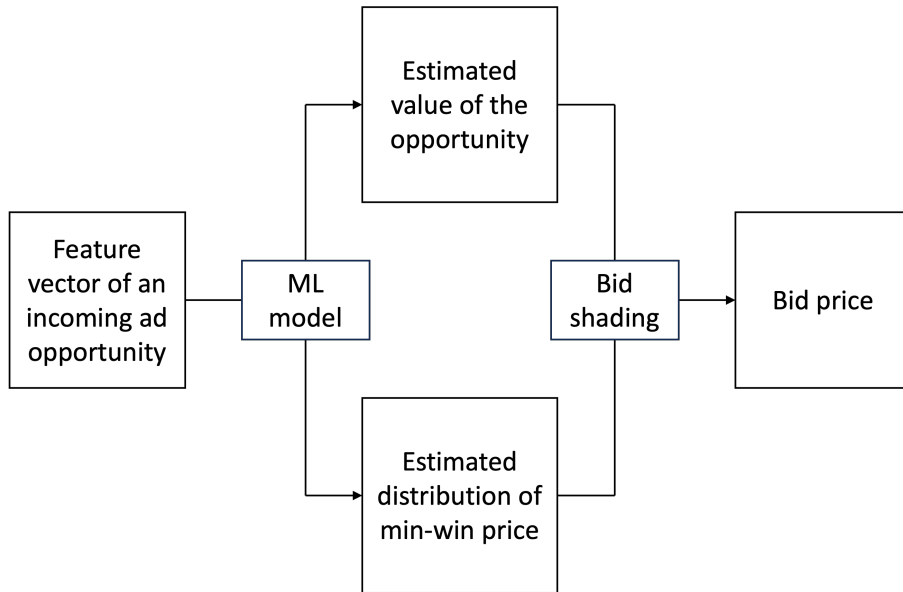
- Problem formulation and computable formula
- Theoretical insight and efficient implementation

3. Experiment

- Dataset, metric, and spend equating
- Outperformance and interpretation

Motivation

- Prediction phase (left)
 - Value & min-win price as output
 - Complex nature + latency constraint = very noisy estimate
- Bid shading phase (right)
 - Value & min-win price as input
 - Robustness against relatively significant estimation errors
 - Distributionally robust optimization (DRO)



Distributionally Robust Bid Shading – Problem Formulation

- Realized value V with distribution \bar{P}
- Min-win price X with distribution \bar{Q}
- Choose bid price b to maximize the expected surplus (baseline)
- Estimates \bar{P} and \bar{Q} tend to be noisy
- Introduce ambiguity sets for them
- Choose bid price b to maximize the worst-case expected surplus (DRBS)

$$\max_b \mathbb{E}_{\bar{P}, \bar{Q}}(V - b)I(X \leq b)$$



$$\max_{b \geq 0} \min_{\substack{P \in \mathcal{P}(\delta_V) \\ Q \in \mathcal{Q}(\delta_X)}} \mathbb{E}_{P, Q}(V - b)I(X \leq b)$$

$$\mathcal{P}(\delta_V) = \{P : D(P || \bar{P}) \leq \delta_V\}$$

$$\mathcal{Q}(\delta_X) = \{Q : D(Q || \bar{Q}) \leq \delta_X\}$$

$$D(p_1 || p_2) = \mathbb{E}_{p_1} \log(p_1(Z)/p_2(X))$$

Distributionally Robust Bid Shading – Computable Formula

The double DRO problem turns out to be almost analytically solvable.

Notations: click probability $\bar{p} = P_{\bar{p}}(V > 0)$, click reward a , value $v = \mathbb{E}V = a\bar{p}$, worst-case value $\bar{v} = ar^{-1}(\delta_V)$, $r(p) = p \log(p/\bar{p}) + (1-p) \log((1-p)/(1-\bar{p}))$, worst-case baseline policy \underline{v} , CDF and PDF of X under \bar{Q} : F and f .

Assumptions: V and X are independent; $\delta_V < r(0)$ and $\delta_X < -\log(1-F(\bar{v}))$; F is log-concave; $F(\bar{v}) < 1/2$ and $F(0) = 0$.

DRBS policy b^* : the unique solution of $g(b) = \delta_X$ in $[\underline{v}, \bar{v}]$ where

$$g(b) = \log \eta(b) - \log J(b) - \frac{F(b) \log \eta(b)}{J(b)},$$

$$J(b) = F(b) + \eta(b) - F(b)\eta(b), \quad \eta(b) = h^{-1}(L(b))$$

$$L(b) = \frac{F(b)}{(\bar{v} - b)f(b)}, \quad h(x) = \frac{x - 1}{\log x}, \quad x \geq 0.$$

Distributionally Robust Bid Shading – Theoretical Insight

DRBS is increasing in δ_X but decreasing in δ_V .

- When we are uncertain about the competitive landscape ($\delta_X > 0$), the competition is fiercer than expected in the worst case, so we should bid higher than the baseline to maintain our win rate.
- When we are uncertain about the value ($\delta_V > 0$), the ad opportunity is less valuable than expected in the worst case, so we should bid lower than the base line to maintain a positive profit margin.
- One KL-ball is not enough. The reduced DRBS either always bids higher or always bids lower.

DRBS bids higher (or lower) than the baseline when v is large (or small) enough.

- When the value is oddly high, why not bid aggressively to secure the deal?
- When the value is oddly low, why not bid conservatively to avoid “winning a loss”?
- Two KL-balls are essential. DRBS with $\delta_X, \delta_V > 0$ can reasonably decide to bid higher or lower.

Distributionally Robust Bid Shading – Efficient Implementation

DRBS policy b^* : the unique solution of $g(b) = \delta_X$ in $[\underline{v}, \bar{v}]$ where

$$g(b) = \log \eta(b) - \log J(b) - \frac{F(b) \log \eta(b)}{J(b)},$$

$$J(b) = F(b) + \eta(b) - F(b)\eta(b), \quad \eta(b) = h^{-1}(L(b))$$

$$L(b) = \frac{F(b)}{(\bar{v} - b)f(b)}, \quad h(x) = \frac{x - 1}{\log x}, \quad x \geq 0.$$

- Both \underline{v} and \bar{v} are easy to compute.
- Since g is strictly increasing in $[\underline{v}, \bar{v}]$, b^* can be computed via bisection.
- For $y > 1$, $h^{-1}(y) = -y \cdot W_{-1}(-e^{-1/y}/y)$, which can be computed via `scipy.special.lambertw`.
- $j_{2/3}(y) \leq h^{-1}(y) \leq j_1(y)$, $j_c(y) = (1 + \sqrt{2} \cdot l(y) + c \cdot l^2(y))$, $l(y) = \sqrt{1/y + \log y - 1}$.

Experiment - Dataset

- Yahoo DSP private bidding dataset on Google Ad Exchange
- Information of 2M bid requests
- More than 1K lines (campaigns)
- Baseline: log-normal model (\bar{Q})
- No private values in public datasets

| Field | Description |
|----------------|--|
| line_id | ID of the line corresponding to the ad that the DSP wants to win the opportunity for |
| ceiling, floor | Range of allowed bid prices |
| mu, sigma | Two estimated lognormal parameters of the distribution of the minimum winning price |
| click_prob | Estimated probability of the ad being clicked |
| click_reward | Reward if the ad is actually clicked |
| value | Product of the above two |
| min_win_price | Actual minimum winning price |

Experiment - Metric

- Value v , bid price b , min-win price X , effective value per dollar spent

$$R = \frac{\sum_i v_i I(X_i \leq b_i)}{\sum_i b_i I(X_i \leq b_i)}$$


- Line k , DRBS R_k^D , baseline R_k^B , percentage improvement

$$\Delta R_k = (R_k^D / R_k^B - 1) \times 100\%$$

- Spend s_k , spend-weighted average over all lines

$$\Delta R = \frac{\sum_k s_k \Delta R_k}{\sum_k s_k}$$



Experiment - Spend Equating



- Spend equating $s_k^D = s_k^B$ is crucial to make offline comparison meaningful in online sense.
- Regulated by the controller (v -modifier), different policies spend the same on average.
- To mimic the controller offline, a standard practice is uniformly modifying v 's for the new policy to make it spend the same as the old one.
- This modification violates the problem formulation where DRBS and baseline are facing the same set of v 's.

Experiment - Spend Equating via Delta Balancing

- $b_i^*(\delta_X, 0) > b_i^*(0,0)$, $b_i^*(0, \delta_V) < b_i^*(0,0)$
- Can we make $b_i^*(\delta_X, \delta_V) = b_i^*(0,0)$ on average?
- First, choose δ_X to maximize the total surplus $\sum_i (v_i - b_i^*(\delta_X, 0)) I(X_i \leq b_i^*(\delta_X, 0))$.
- Second, choose δ_V to equate the spend $\sum_i b_i^*(\delta_X, \delta_V) = \sum_i b_i^*(0,0)$.
- The resulting DRBS policy handles two sources of uncertainty while maintaining the same spend rate as the baseline policy.
- Two KL-balls are essential.



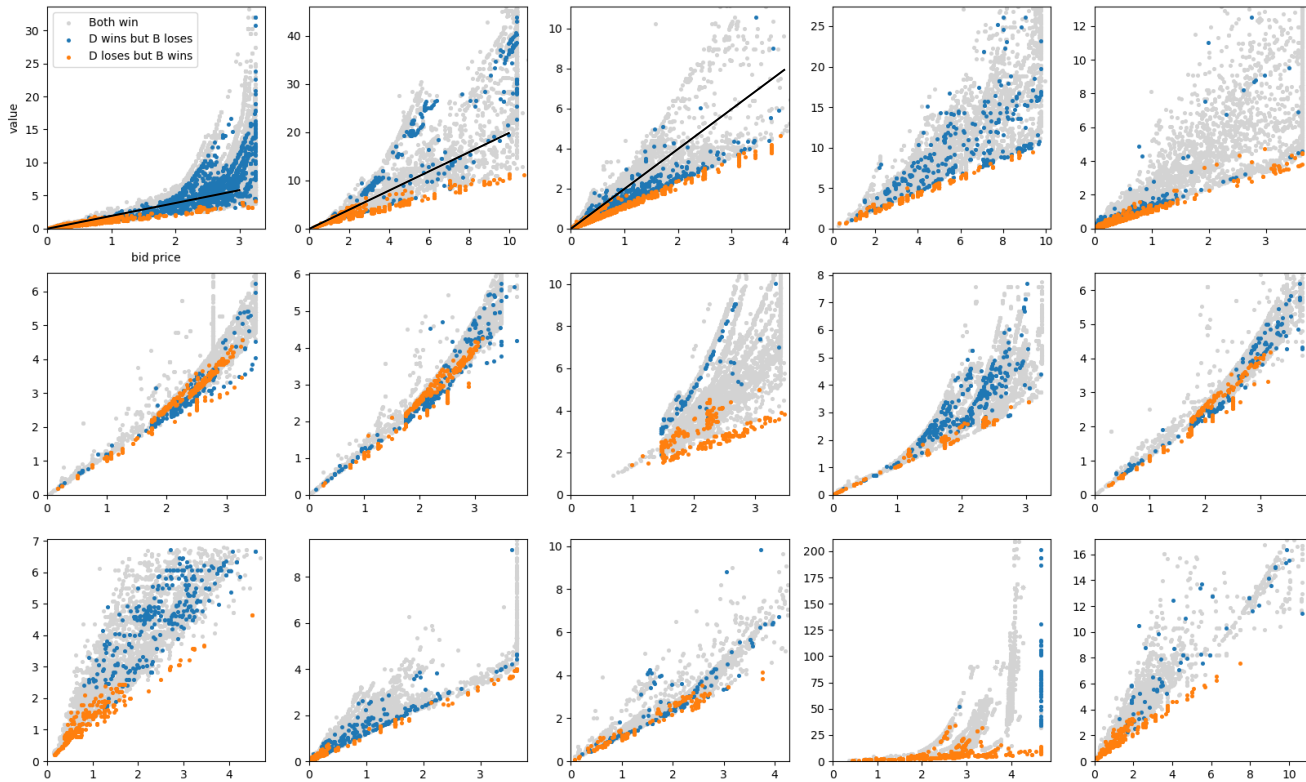
Experiment - Outperformance and Interpretation

- For each line, we use 25% of the data to compute δ_X and δ_V while the rest 75% is used for testing.
- For the 3 largest lines with spend weights 8.6%, 8.6%, 6.6%, $\Delta R_1 = 1.0\%$, $\Delta R_2 = 2.3\%$, $\Delta R_3 = 0.0\%$.
- For all lines, the spend-weighted average $\Delta R = 0.65\%$. Where does the gain come from?
- Exchange low- v/b wins for high- v/b wins.

| X | b^B | b^D | v | v/b |
|------|-------|-------|-------|-------|
| 2.47 | 2.29 | 2.49* | 6.64 | 2.67 |
| 0.59 | 0.58 | 0.60* | 0.89 | 1.48 |
| 2.09 | 2.00 | 2.20* | 6.83 | 3.10 |
| 2.96 | 2.93 | 3.11* | 10.25 | 3.30 |
| 1.78 | 1.77 | 1.86* | 4.22 | 2.27 |

| X | b^B | b^D | v | v/b |
|------|-------|-------|------|-------|
| 0.13 | 0.13* | 0.00 | 0.25 | 1.92 |
| 0.13 | 0.13* | 0.00 | 0.25 | 1.92 |
| 0.34 | 0.34* | 0.00 | 0.40 | 1.18 |
| 0.32 | 0.33* | 0.00 | 0.43 | 1.30 |
| 0.24 | 0.24* | 0.24 | 0.39 | 1.63 |

Experiment - Outperformance and Interpretation



Takeaway

- Real-time bidding algorithm needs to be robust.
- Double distributionally robust optimization works.
- Two KL-balls are essential (for spend equating).
- DRBS policy has computable formula and interpretable outperformance.

Thank You

<https://quyanlin.github.io>

